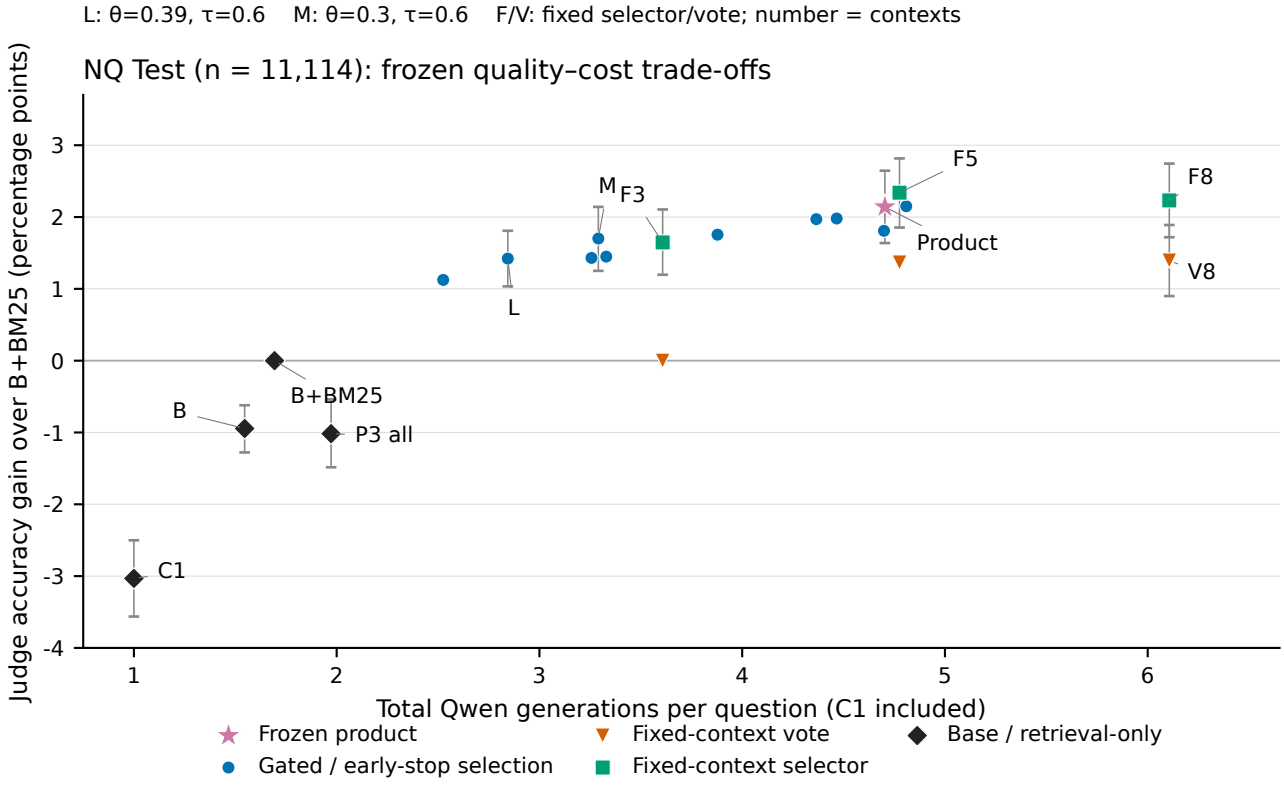
L: θ=0.39, τ=0.6 M: θ=0.3, τ=0.6 F/V: fixed selector/vote; number = contexts
NQ Test (n = 11,114): frozen quality–cost trade-offs
Judge accuracy gain over B+BM25 (percentage points)
Total Qwen generations per question (C1 included)
C1
B
B+BM25
P3 all
L
M
F3
F5
Product
F8
V8
Frozen product
Gated / early-stop selection
Fixed-context vote
Fixed-context selector
Base / retrieval-only
Whiskers: paired item bootstrap 95% CIs (10,000 resamples). All points are existing policies; no new selection.